

▼ NOTEBOOK TO DEVELOP AN IMAGE CAPTIONING MODEL

This notebook is intended to develop a model which will take in as input an image and output the captions from it, that is to say, it will describe the different objects present in the image.

MOUNTING GDRIVE ONTO THE NOTEBOOK

```
from google.colab import drive
drive.mount('/content/drive', force_remount=True)
COLAB = True
```

Mounted at /content/drive

IMPORTING NECESSARY LIBRARIES

```
import os
import string
import glob
from tensorflow.keras.applications import MobileNet
import tensorflow.keras.applications.mobilenet

from tensorflow.keras.applications.inception_v3 import InceptionV3
import tensorflow.keras.applications.inception_v3

from tqdm import tqdm
import tensorflow.keras.preprocessing.image
import pickle
from time import time
import numpy as np
from PIL import Image
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import (LSTM, Embedding,
    TimeDistributed, Dense, RepeatVector,
    Activation, Flatten, Reshape, concatenate,
    Dropout, BatchNormalization)
from tensorflow.keras.optimizers import Adam, RMSprop
from tensorflow.keras import Input, layers
from tensorflow.keras import optimizers

from tensorflow.keras.models import Model

from tensorflow.keras.layers import add
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.utils import to_categorical
import matplotlib.pyplot as plt

START = "startseq"
STOP = "endseq"
```

```
EPOCHS = 10
USE_INCEPTION = True
```

```
# Nicely formatted time string
def hms_string(sec_elapsed):
    h = int(sec_elapsed / (60 * 60))
    m = int((sec_elapsed % (60 * 60)) / 60)
    s = sec_elapsed % 60
    return f"{h}:{m:>02}:{s:>05.2f}"
```

ROOT_CAPTIONING WILL HOLD THE ROOT ADDRESS OF ALL THE FILES AND FOLDERS BEING USED IN THIS NOTEBOOK

▼ DATASETS NEEDED

For the development of this project the Flickr8k dataset has been used to train the images against captions. It is an open-source dataset containing 8k images each with at least one object. Corresponding to each image we have a captions that accurately describes the image. The captions are stored in the Folder Flickr8k_text->Flickr8k.token.txt. This tokens data contains up to 5 captions for each image [resent in the Flickr8k_image dataset.

Along with flickr8k dataset we also used the glove.6B dataset which stores the vector representation of common words we encounter in our day-to-day life.

```
root_captioning = "/content/drive/MyDrive/Colab Notebooks/CAPTIONING"
```

▼ PREPARING THE DATASET

Each image in the dataset contains up to 5 captions associated with it. These captions are strings which may have unnecessary punctuations and/or spaces in between. We must clean these irregularities. We also prepare a dictionary, lookup, where each entry contains the list of all the image caption captions for a particular image index. The lex set contains all the unique words present in the tokens dataset.

```
null_punct = str.maketrans('', '', string.punctuation)
lookup = dict()

with open( os.path.join(root_captioning, 'Flickr8k_text', 'Flickr8k.token.txt'), 'r') as fp:

    max_length = 0
    for line in fp.read().split('\n'):
        tok = line.split()
        if len(line) >= 2:
            id = tok[0].split('.')[0]
            desc = tok[1:]
```

```
# Cleanup description
desc = [word.lower() for word in desc]
desc = [w.translate(null_punct) for w in desc]
desc = [word for word in desc if len(word)>1]
desc = [word for word in desc if word.isalpha()]
max_length = max(max_length, len(desc))

if id not in lookup:
    lookup[id] = list()
lookup[id].append(' '.join(desc))

lex = set()
for key in lookup:
    [lex.update(d.split()) for d in lookup[key]]

print(len(lookup)) # The length of the dictionary
print(len(lex)) # How many unique words
print(max_length) # Maximum length of a caption (in words)
```

```
8092
8763
32
```

▼ FINDING PATH TO THE IMAGES

the glob module defined below is used to find path to all the images present in the Flickr8k_Dataset

```
img = glob.glob(os.path.join(root_captioning, 'Flickr8k_Dataset', '*.jpg'))

print(len(img))
img[0:5]#printing the addresses of the first 5 images

8166
['/content/drive/MyDrive/Colab Notebooks/CAPTIONING/Flickr8k_Dataset/3282634762_2656
'/content/drive/MyDrive/Colab Notebooks/CAPTIONING/Flickr8k_Dataset/3273403495_fcd6
'/content/drive/MyDrive/Colab Notebooks/CAPTIONING/Flickr8k_Dataset/2483993772_f64e
'/content/drive/MyDrive/Colab Notebooks/CAPTIONING/Flickr8k_Dataset/3563461991_de0f
'/content/drive/MyDrive/Colab Notebooks/CAPTIONING/Flickr8k_Dataset/1020651753_0607

train_images_path = os.path.join(root_captioning, 'Flickr8k_text', 'Flickr_8k.trainImages.txt')
train_images = set(open(train_images_path, 'r').read().strip().split('\n'))#STORES THE ADD
test_images_path = os.path.join(root_captioning, 'Flickr8k_text', 'Flickr_8k.testImages.txt')
test_images = set(open(test_images_path, 'r').read().strip().split('\n')) #STORES THE ADDR

train_img = [] #stores the train addresses in a list
test_img = [] #stores the test addresses in a list

for i in img:
```

```

r = os.path.split(1)[-1]
if f in train_images:
    train_img.append(f)
elif f in test_images:
    test_img.append(f)

print(len(train_images))#stores the addresses of all the train images
print(len(test_images))#stores the addresses of all the test images

6000
1000

```

BUILDING A DICTIONARY FOR STORING THE TRAIN IMAGES ALONG WITH ITS CAPTIONS. WE ALSO ADD A SPECIAL START AND END TOKEN IN THE BEGINNING/END SO AS TO KNOW WHEN TO START A CAPTION AND WHEN TO FINISH IT

```

train_descriptions = {k:v for k,v in lookup.items() if f'{k}.jpg' in train_images}
for n,v in train_descriptions.items():
    for d in range(len(v)):
        v[d] = f'{START} {v[d]} {STOP}'

print(len(train_descriptions))

6000

```

▼ EXTRACTING FEATURES FROM IMAGES

THE IMAGES TO BE PASSED TO THE CAPTIONING MODEL CANNOT BE RAW IMAGES, WE HAVE TO EXTRACT THE FEATURES FROM EACH OF THESE IMAGES AN THEN PASS IT ONTO THE CAPTIONING MODEL.FOR DOING THIS WE USE THE PRE-TRAINED INCEPTIONV3 CNN MODEL. THIS FORMS OUR ENCODING MODEL.WITHOUT THE USE OF SUCH PRE-BUILD MODELS THE IMAGES WILL TAKE IN SIGNIFICANTLY MORE TRAINING.

```

encode_model = InceptionV3(weights='imagenet')
encode_model = Model(encode_model.input, encode_model.layers[-2].output)
WIDTH = 299
HEIGHT = 299
OUTPUT_DIM = 2048
preprocess_input = tensorflow.keras.applications.inception_v3.preprocess_input

encode_model.summary()

```

activation_278 (Activation)	(None, 8, 8, 384)	0	batch_normalizati
conv2d_275 (Conv2D)	(None, 8, 8, 384)	442368	activation_274[0]
conv2d_276 (Conv2D)	(None, 8, 8, 384)	442368	activation_274[0]

conv2d_279 (Conv2D)	(None, 8, 8, 384)	442368	activation_278[0]
conv2d_280 (Conv2D)	(None, 8, 8, 384)	442368	activation_278[0]
average_pooling2d_26 (AveragePo	(None, 8, 8, 2048)	0	mixed9[0][0]
conv2d_273 (Conv2D)	(None, 8, 8, 320)	655360	mixed9[0][0]
batch_normalization_275 (BatchN	(None, 8, 8, 384)	1152	conv2d_275[0][0]
batch_normalization_276 (BatchN	(None, 8, 8, 384)	1152	conv2d_276[0][0]
batch_normalization_279 (BatchN	(None, 8, 8, 384)	1152	conv2d_279[0][0]
batch_normalization_280 (BatchN	(None, 8, 8, 384)	1152	conv2d_280[0][0]
conv2d_281 (Conv2D)	(None, 8, 8, 192)	393216	average_pooling2d
batch_normalization_273 (BatchN	(None, 8, 8, 320)	960	conv2d_273[0][0]
activation_275 (Activation)	(None, 8, 8, 384)	0	batch_normalizati
activation_276 (Activation)	(None, 8, 8, 384)	0	batch_normalizati
activation_279 (Activation)	(None, 8, 8, 384)	0	batch_normalizati
activation_280 (Activation)	(None, 8, 8, 384)	0	batch_normalizati
batch_normalization_281 (BatchN	(None, 8, 8, 192)	576	conv2d_281[0][0]
activation_273 (Activation)	(None, 8, 8, 320)	0	batch_normalizati
mixed9_1 (Concatenate)	(None, 8, 8, 768)	0	activation_275[0] activation_276[0]
concatenate_5 (Concatenate)	(None, 8, 8, 768)	0	activation_279[0] activation_280[0]
activation_281 (Activation)	(None, 8, 8, 192)	0	batch_normalizati
mixed10 (Concatenate)	(None, 8, 8, 2048)	0	activation_273[0] mixed9_1[0][0] concatenate_5[0][activation_281[0]
avg_pool (GlobalAveragePooling2	(None, 2048)	0	mixed10[0][0]
=====			
Total params: 21,802,784			
Trainable params: 21,768,352			
Non-trainable params: 34,432			

THIS FUNCTION TAKES IN A RAW IMAGE AND OUTPUTS THE FEATURES EXTRACTED FROM THEM, SO AS TO BE USED FOR LATER MODELS

```
def encodeImage(img):
    img = img.resize((WIDTH, HEIGHT), Image.ANTIALIAS)
```

```

x = tensorflow.keras.preprocessing.image.img_to_array(img)
x = np.expand_dims(x, axis=0)
x = preprocess_input(x)
x = encode_model.predict(x)
x = np.reshape(x, OUTPUT_DIM )
return x

```

▼ STORING THE FEATURES FOR THE TRAIN IMAGES

WE USE THE PICKLE LIBRARY TO STORE THE TRAINING IMAGES ID ALONG WITH ITS EXTRACTED FEATURES IN THE FORM OF A DICTIONARY. DOING THIS WILL SAVE US A LOT OF TIME LATER AND ALSO PREVENT US FROM RE-EXTRACTING THE SAME FEATURES IF A SITUATION OF RE-TRAINING OCCURS

```

train_path = os.path.join(root_captioning, "data", f'train{OUTPUT_DIM}.pkl')
if not os.path.exists(train_path):
    start = time()
    encoding_train = {}
    for id in tqdm(train_img): #iterating the traing images by using the tqdm function so th
        image_path = os.path.join(root_captioning, 'Flicker8k_Dataset', id)
        img = tensorflow.keras.preprocessing.image.load_img(image_path, target_size=(HEIGHT, WI
        encoding_train[id] = encodeImage(img)
    with open(train_path, "wb") as fp:
        pickle.dump(encoding_train, fp)
    print(f"\nGenerating training set took: {hms_string(time()-start)}")
else:
    with open(train_path, "rb") as fp:
        encoding_train = pickle.load(fp)

```

▼ STORING FEATURES FOR TEST IMAGES

SAME AS WE DID FOR THE TRAIN IMAGES, WE ALSO STORE THE FEATURES FOR THE TEST IMAGES INTO PICKLE FILE FOR LATER USE

```

test_path = os.path.join(root_captioning, "data", f'test{OUTPUT_DIM}.pkl')
if not os.path.exists(test_path):
    start = time()
    encoding_test = {}
    for id in tqdm(test_img):
        image_path = os.path.join(root_captioning, 'Flicker8k_Dataset', id)
        img = tensorflow.keras.preprocessing.image.load_img(image_path, target_size=(HEIGHT, WI
        encoding_test[id] = encodeImage(img)
    with open(test_path, "wb") as fp:
        pickle.dump(encoding_test, fp)
    print(f"\nGenerating testing set took: {hms_string(time()-start)}")
else:
    with open(test_path, "rb") as fp:
        encoding_test = pickle.load(fp)

```

WE SEPRATE CAPTIONS FROM THE IMAGES, THERE ARE TWO SIDE OF TRAINING THE MODEL,ONE ON IMAGES AND THE OTHER ON CAPTIONS

```
all_train_captions = []
for key, val in train_descriptions.items():
    for cap in val:
        all_train_captions.append(cap)
len(all_train_captions)
```

30000

LESS COMMON WORDS CAN BE BE MISLEADING TO THE NEURAL NETWORK BEING TRAINED, SO WE REMOVE ANY WORD WHICH OCCURS LESS THAN 10 TIMES AND PREPARE A RESULTING VOCABULARY

```
word_count_threshold = 10
word_counts = {} #STORES THE NUMBER OF TIMES EACH WORD HAS OCCURED IN THE CAPTIONS TRAININ
nsents = 0
for sent in all_train_captions:
    nsents += 1
    for w in sent.split(' '):
        word_counts[w] = word_counts.get(w, 0) + 1
```

```
vocab = [w for w in word_counts if word_counts[w] >= word_count_threshold]# VOCAB STORES 0
print('preprocessed words %d ==> %d' % (len(word_counts), len(vocab)))
```

preprocessed words 7578 ==> 1651

NOW WE PREPARE TWO DICTIONARIES FOR THE VOCAB, THE IDXTOWORD DICT ASSOCIATES THE INDEX TO CORRESPONDING WORDS PRESENT IN THE VOCAB LIST . TH WORDTIDX DOES THE OPPOSITE.

```
idxtoword = {}
wordtoidx = {}

ix = 1
for w in vocab:
    wordtoidx[w] = ix
    idxtoword[ix] = w
    ix += 1
```

```
vocab_size = len(idxtoword) + 1
vocab_size
```

1652

▼ STORING THE DICTIONARIES

WE NOW STORE THE IDXTOWORD AND THE WORDTOIDX DICTIONARIES INTO THE GDRIVE IN THE FORM OF PICKLE FILES.THESE WILL BE USED LATER FOR PREDICTING THE RESULTS.

```
idxtoword_path=os.path.join(root_captioning,"data",'idxtoword.pkl')
with open(idxtoword_path, "wb") as fp:
    pickle.dump(idxtoword, fp)
print("data saved")
```

data saved

```
wordtoidx_path=os.path.join(root_captioning,"data",'wordtoidx.pkl')
with open(wordtoidx_path, "wb") as fp:
    pickle.dump(wordtoidx,fp)
print("data saved")
```

data saved

```
max_length +=2
print(max_length)
```

34

▼ USING DATA GENERATORS

Up to this point, we've always generated training data ahead of time and fit the neural network to it. It is not always practical to create all of the training data ahead of time. The memory demands can be considerable. If we generate the training data as the neural network needs it, it is possible to use a Keras generator. The generator will create new data as it is needed. The generator provided here creates the training data for the caption neural network, as it is needed.

```
def data_generator(descriptions, photos, wordtoidx, \
                  max_length, num_photos_per_batch):
    # x1 - Training data for photos
    # x2 - The caption that goes with each photo
    # y - The predicted rest of the caption
    x1, x2, y = [], [], []
    n=0
    while True:
        for key, desc_list in descriptions.items():
            n+=1
            photo = photos[key+'.jpg']
            # Each photo has 5 descriptions
            for desc in desc_list:
                # Convert each caption into a list of sequences.
                seq = [wordtoidx[word] for word in desc.split(' ') if word in wordtoidx]
```



```
# Generate a training case for every possible sequence and outcome
for i in range(1, len(seq)):
    in_seq, out_seq = seq[:i], seq[i]
    in_seq = pad_sequences([in_seq], maxlen=max_length)[0]
    out_seq = to_categorical([out_seq], num_classes=vocab_size)[0]
    x1.append(photo)
    x2.append(in_seq)
    y.append(out_seq)
if n==num_photos_per_batch:
    yield ([np.array(x1), np.array(x2)], np.array(y))
    x1, x2, y = [], [], []
    n=0
```

GLOVE IS AN UNSUPERVISED LEARNING ALGORITHM FOR OBTAINING THE VECTOR REPRESENTATION OF THE WORDS. THESE VECTORS ARE GLOBAL IN NATURE AND ARE USED FOR TRAININGS MODELS WHICH TAKE IN WORDS(TEXT) AS INPUT.

```
glove_dir = os.path.join(root_captioning, 'glove.6B')
embeddings_index = {}
f = open(os.path.join(glove_dir, 'glove.6B.200d.txt'), encoding="utf-8")

for line in tqdm(f):
    values = line.split()
    word = values[0]
    coefs = np.asarray(values[1:], dtype='float32')
    embeddings_index[word] = coefs

f.close()
print(f'Found {len(embeddings_index)} word vectors.')
```

```
400000it [00:23, 16940.93it/s]Found 400000 word vectors.
```

▼ CREATING EMBEDDING MATRIX

NOW WE USE THE GLOVE EMBEDDINGS TO CREATE A 200 DIM VECTOR FOR EACH WORD PRESENT IN OUR VOCABULARY. WE STORE THESE IN A MATRIX CALLED EMBEDDING MATRIX

```
embedding_dim = 200

# Get 200-dim dense vector for each word in our vocabulary
embedding_matrix = np.zeros((vocab_size, embedding_dim))

for word, i in wordtoidx.items():
    #if i < max_words:
    embedding_vector = embeddings_index.get(word)
    if embedding_vector is not None:
        # Words not found in the embedding index will be all zeros
        embedding_matrix[i] = embedding_vector
```

```
embedding_matrix.shape
```

```
(1652, 200)
```

▼ CREATING THE CAPTION MODEL

NEXT WE HAVE TO CREATE THE CAPTION MODEL, THIS MODEL TAKES IN AS INPUT AN IMAGE AND OUTPUTS THE PREDICTED CAPTION

```
inputs1 = Input(shape=(OUTPUT_DIM,))
fe1 = Dropout(0.5)(inputs1)
fe2 = Dense(256, activation='relu')(fe1)
inputs2 = Input(shape=(max_length,))
se1 = Embedding(vocab_size, embedding_dim, mask_zero=True)(inputs2)
se2 = Dropout(0.5)(se1)
se3 = LSTM(256)(se2)
decoder1 = add([fe2, se3])
decoder2 = Dense(256, activation='relu')(decoder1)
outputs = Dense(vocab_size, activation='softmax')(decoder2)
caption_model = Model(inputs=[inputs1, inputs2], outputs=outputs)
```

```
embedding_dim
```

```
200
```

```
caption_model.summary()
```

```
Model: "model_4"
```

Layer (type)	Output Shape	Param #	Connected to
input_7 (InputLayer)	[(None, 34)]	0	
input_6 (InputLayer)	[(None, 2048)]	0	
embedding_1 (Embedding)	(None, 34, 200)	330400	input_7[0][0]
dropout_2 (Dropout)	(None, 2048)	0	input_6[0][0]
dropout_3 (Dropout)	(None, 34, 200)	0	embedding_1[0][0]
dense_3 (Dense)	(None, 256)	524544	dropout_2[0][0]
lstm_1 (LSTM)	(None, 256)	467968	dropout_3[0][0]
add_1 (Add)	(None, 256)	0	dense_3[0][0] lstm_1[0][0]
dense_4 (Dense)	(None, 256)	65792	add_1[0][0]
dense_5 (Dense)	(None, 1652)	424564	dense_4[0][0]
Total params: 1,813,268			

Trainable params: 1,813,268
 Non-trainable params: 0

```
caption_model.layers[2].set_weights([embedding_matrix])
caption_model.layers[2].trainable = False
caption_model.compile(loss='categorical_crossentropy', optimizer='adam')
```

```
number_pics_per_batch = 3
steps = len(train_descriptions)//number_pics_per_batch
```

▼ TRAINING THE CAPTION MODEL

```
model_path = os.path.join(root_captioning, "data", f'caption-model.hdf5')
if not os.path.exists(model_path):
    for i in tqdm(range(EPOCHS*2)):
        generator = data_generator(train_descriptions, encoding_train,
                                   wordtoidx, max_length, number_pics_per_batch)
        caption_model.fit_generator(generator, epochs=1,
                                    steps_per_epoch=steps, verbose=1)

    caption_model.optimizer.lr = 1e-4
    number_pics_per_batch = 6
    steps = len(train_descriptions)//number_pics_per_batch

    for i in range(EPOCHS):
        generator = data_generator(train_descriptions, encoding_train,
                                   wordtoidx, max_length, number_pics_per_batch)
        caption_model.fit_generator(generator, epochs=1,
                                    steps_per_epoch=steps, verbose=1)
    caption_model.save_weights(model_path)
    print(f"\nTraining took: {hms_string(time()-start)}")
else:
    caption_model.load_weights(model_path)
```

▼ FUNCTION TO OUTPUT THE RESULT

THE GENERATECAPTION FUNCTION IS USED TO TAKE A ENCODED IMAGE AND OUTPUT THE PREDICTED CAPTIONS.

```
def generateCaption(photo):
    in_text = START
    for i in range(max_length):
        sequence = [wordtoidx[w] for w in in_text.split() if w in wordtoidx]
```

```

sequence = pad_sequences([sequence], maxlen=max_length)
yhat = caption_model.predict([photo,sequence], verbose=0)
yhat = np.argmax(yhat)
word = idxtoword[yhat]
in_text += ' ' + word
if word == STOP:
    break
final = in_text.split()
final = final[1:-1]
final = ' '.join(final)
return final

```

▼ TESTING THE DATA

NOW WE TEST THE MODEL AGAINST A FEW OF THE TRAIN IMAGES

```

for z in range(2): # set higher to see more examples
    pic = list(encoding_test.keys())[z]
    image = encoding_test[pic].reshape((1,OUTPUT_DIM))
    print(os.path.join(root_captioning,'Flicker8k_Dataset', pic))
    x=plt.imread(os.path.join(root_captioning,'Flicker8k_Dataset', pic))
    plt.imshow(x)
    plt.show()
    print("Caption:",generateCaption(image))
    print("_____")


```

/content/drive/MyDrive/Colab Notebooks/CAPTIONING/Flicker8k_Dataset/3192266178_f9bf5c



encoding_test[pic].shape

(2048,)



Caption: man in blue shirt is playing with dog in the grass

/content/drive/MyDrive/Colab Notebooks/CAPTIONING/Flicker8k_Dataset/3234115903_f4dfc8



Caption: two men in martial arts match

✓ 0s completed at 9:42 PM